

B.M.S. COLLEGE OF ENGINEERING BENGALURU

Autonomous Institute, Affiliated to VTU



Lab Record

BIG DATA ANALYTICS

*Submitted in partial fulfillment for the 6th Semester
Laboratory*

Bachelor of Technology
in
Computer Science and Engineering

Submitted by:

SHREEHARI KULKARNI
1BM19CS153

Department of Computer Science and Engineering
B.M.S. College of Engineering
Bull Temple Road, Basavanagudi, Bangalore 560
019 Mar-June 2021

B.M.S. COLLEGE OF ENGINEERING
DEPARTMENT OF COMPUTER SCIENCE AND
ENGINEERING



CERTIFICATE

This is to certify that the Lab work entitled “BIG DATA ANALYTICS” carried out SHREEHARI KULKARNI(**1BM19CS153**), who is a bonafide student of B. M. S. College of Engineering. It is in partial fulfillment for the award of Bachelor of Engineering in Computer Science and Engineering of the Visvesvaraya Technological University, Belgaum during the year 2022. The Lab report has been approved as it satisfies the academic requirements in respect of a **BIG DATA ANALYTICS- (20CS6PEBDA)** work prescribed for the said degree.

.

ANTARA ROY CHOUDHURY
Assistant Professor,
Department of CSE,
BMSCE, Bengaluru.

Dr. Jyothi S Nayak,
Professor and Head,
Department of CSE,
BMSCE, Bengaluru.

TABLE OF CONTENTS

SL NO	TITLE
1	EMPLOYEE DATABASE
2	LIBRARY DATABASE
3	MONGODB SAMPLE
4	HADOOP INSTALLATION
5	HADOOP SAMPLE
6	MAPREDUCE TEMPERATURE
7	MAPREDUCE TOPN
8	MAPREDUCE JOIN
9	SCALA HELLO WORD
10	SPARK RDD AND FLAT MAP

Date - 19/04/2022

Employee database (CASSANDRA)

Question -

Perform the following DB operations using Cassandra.

1. Create a keyspace by name Employee
2. Create a column family by name Employee-Info with attributes Emp_Id Primary Key, Emp_Name, Designation, Date_of_Joining, Salary, Dept_Name
3. Insert the values into the table in batch
3. Update Employee name and Department of Emp-Id 121
4. Sort the details of Employee records based on salary
5. Alter the schema of the table Employee_Info to add a column Projects which stores a set of Projects done by the corresponding Employee.
6. Update the altered table to add project names.
7. Create a TTL of 15 seconds to display the values of Employees.

```
cqlsh> create keyspace employee_info with  
replication={ 'class': 'SimpleStrategy', 'replication_factor': 1 };
```

```
cqlsh> use employee_info;
```

```
cqlsh:employee_info> create table employee_details(emp_id int, emp_name text, designation  
text, doj timestamp, salary double, dept_name text, primary key(emp_id, salary));
```

```
cqlsh:employee_info> begin batch  
... insert into  
employee_details(emp_id, emp_name, designation, doj, salary, dept_name) values  
(100, 'tanya', 'manager', '2020-09-11', 30000, 'testing')  
... insert into  
employee_details(emp_id, emp_name, designation, doj, salary, dept_name) values
```

```
(111,'sriram','associate','2020-06-11',25000,'development')
```

```
... insert into
```

```
employee_details(emp_id,emp_name,designation,doj,salary,dept_name) values
```

```
(121,'shiva','manager','2020-01-03',35000,'hr')
```

... apply batch;

cqlsh:employee_info> select * from employee_details;

emp_id	salary	dept_name	designation	doj	emp_name
_____	_____	_____	_____	_____	_____
111	25000	development	associate	2020-06-10 18:30:00.000000+0000	sriram
121	35000	hr	manager	2020-01-02 18:30:00.000000+0000	shiva
100	30000	testing	manager	2020-09-10 18:30:00.000000+0000	tanya

(3 rows)

cqlsh:employee_info> update employee_details set emp_name='shaan' where emp_id=121 and salary=35000;

cqlsh:employee_info> select * from employee_details;

emp_id	salary	dept_name	designation	doj	emp_name
_____	_____	_____	_____	_____	_____
111	25000	development	associate	2020-06-10 18:30:00.000000+0000	sriram
121	35000	hr	manager	2020-01-02 18:30:00.000000+0000	shaan
100	30000	testing	manager	2020-09-10 18:30:00.000000+0000	tanya

(3 rows)

cqlsh:employee_info> alter table employee_details add project text;

cqlsh:employee_info> update employee_details set project='chat app' where emp_id=111 and salary=25000;

```
cqlsh:employee_info> update employee_details set project='campusx' where emp_id=121
and salary=35000;
```

```
cqlsh:employee_info> update employee_details set project='canteen app' where emp_id=100 and
salary=30000;
```

```
cqlsh:employee_info> select * from employee_details;
```

emp_id	salary	dept_name	designation	doj	emp_name	project
111	25000	development	associate	2020-06-10 18:30:00.000000+0000	sriram	chat app
121	35000	hr	manager	2020-01-02 18:30:00.000000+0000	shaan	
campusx 100	30000	testing			manager	
2020-09-10 18:30:00.000000+0000					tanya	canteen app

(3 rows)

```
cqlsh:employee_info> insert into
employee_details(emp_id,emp_name,designation,doj,salary,dept_name)
values(113,'sam','manager','2020-09-09',30000,'testing') using ttl 30;
```

```
cqlsh:employee_info> select ttl(emp_name) from employee_details where emp_id=113 and
salary=30000;
```

ttl(emp_name)

29

(1 rows)

```
cqlsh:employee_info> select * from employee_details;
```

emp_id	salary	dept_name	designation	doj	emp_name	project
_____	_____	_____	_____	_____	_____	_____
111	25000	development	associate	2020-06-10 18:30:00.000000+0000	sriram	chat app
113	30000	testing	manager	2020-09-08 18:30:00.000000+0000	sam	null
121	35000	hr	manager	2020-01-02 18:30:00.000000+0000	shaan	
campusx 100	30000	testing			manager	
2020-09-10 18:30:00.000000+0000					tanya	canteen

app

(4 rows)

```
cqlsh:employee_info> select * from employee_details;
```

emp_id	salary	dept_name	designation	doj	emp_name	project
_____	_____	_____	_____	_____	_____	_____
111	25000	development	associate	2020-06-10 18:30:00.000000+0000	sriram	chat app
121	35000	hr	manager	2020-01-02 18:30:00.000000+0000	shaan	
campusx 100	30000	testing			manager	
2020-09-10 18:30:00.000000+0000					tanya	canteen

app

(3 rows)

```
cqlsh:employee_info> paging off;
```

Disabled Query paging.

```
cqlsh:employee_info> select * from employee_details where emp_id in (111,121,100) order by salary;
```


emp_id	salary	dept_name	designation	doj	emp_name	project
111	25000	development	associate	2020-06-10 18:30:00.000000+0000	sriram	chat app
100	30000	testing	manager	2020-09-10 18:30:00.000000+0000	tanya	canteen app
121	35000	hr	manager	2020-01-02 18:30:00.000000+0000	shaan	campusx

(3 rows)

SCREENSHOTS -

```

cqlsh> create keyspace employee_info with replication={'class':'SimpleStrategy','replication_factor':1};
cqlsh> use employee_info;

```

```

cqlsh:employee_info> create table employee_details(emp_id int,emp_name text,designation text,doj timestamp,salary double,dept_name text,primary key(emp_id,salary));
cqlsh:employee_info> begin batch
... insert into employee_details(emp_id,emp_name,designation,doj,salary,dept_name) values (100,'tanya','manager','2020-09-11',30000,'testing')
... insert into employee_details(emp_id,emp_name,designation,doj,salary,dept_name) values (111,'sriram','associate','2020-06-11',25000,'development')
... insert into employee_details(emp_id,emp_name,designation,doj,salary,dept_name) values (121,'shiva','manager','2020-01-03',35000,'hr')
... apply batch;
cqlsh:employee_info> select * from employee_details;

```

emp_id	salary	dept_name	designation	doj	emp_name
111	25000	development	associate	2020-06-10 18:30:00.000000+0000	sriram
121	35000	hr	manager	2020-01-02 18:30:00.000000+0000	shiva
100	30000	testing	manager	2020-09-10 18:30:00.000000+0000	tanya

(3 rows)

```

cqlsh:employee_info> update employee_details set emp_name='shaan' where emp_id=121 and salary=35000;
cqlsh:employee_info> select * from employee_details;

```

emp_id	salary	dept_name	designation	doj	emp_name
111	25000	development	associate	2020-06-10 18:30:00.000000+0000	sriram
121	35000	hr	manager	2020-01-02 18:30:00.000000+0000	shaan
100	30000	testing	manager	2020-09-10 18:30:00.000000+0000	tanya

(3 rows)

```

cqlsh:employee_info> alter table employee_details add project text;

```

```
cqlsh:employee_info> update employee_details set project='chat app' where emp_id=111 and salary=25000;
cqlsh:employee_info> update employee_details set project='campusx' where emp_id=121 and salary=35000;
cqlsh:employee_info> update employee_details set project='canteen app' where emp_id=100 and salary=30000;
cqlsh:employee_info> select * from employee_details;
```

emp_id	salary	dept_name	designation	doj	emp_name	project
111	25000	development	associate	2020-06-10 18:30:00.000000+0000	sriram	chat app
121	35000	hr	manager	2020-01-02 18:30:00.000000+0000	shaan	campusx
100	30000	testing	manager	2020-09-10 18:30:00.000000+0000	tanya	canteen app

(3 rows)

```
cqlsh:employee_info> insert into employee_details(emp_id,emp_name,designation,doj,salary,dept_name) values(113,'sam','manager','2020-09-09',30000,'testing') using ttl 30;
```

```
cqlsh:employee_info> select ttl(emp_name) from employee_details where emp_id=113 and salary=30000;
```

```
ttl(emp_name)
-----
```

(0 rows)

```
cqlsh:employee_info> insert into employee_details(emp_id,emp_name,designation,doj,salary,dept_name) values(113,'sam','manager','2020-09-09',30000,'testing') using ttl 30;
```

```
cqlsh:employee_info> select ttl(emp_name) from employee_details where emp_id=113 and salary=30000;
```

```
ttl(emp_name)
-----
```

29

(1 rows)

```
cqlsh:employee_info> select * from employee_details;
```

emp_id	salary	dept_name	designation	doj	emp_name	project
111	25000	development	associate	2020-06-10 18:30:00.000000+0000	sriram	chat app
113	30000	testing	manager	2020-09-08 18:30:00.000000+0000	sam	null
121	35000	hr	manager	2020-01-02 18:30:00.000000+0000	shaan	campusx
100	30000	testing	manager	2020-09-10 18:30:00.000000+0000	tanya	canteen app

(4 rows)

```
cqlsh:employee_info> select * from employee_details;
```

emp_id	salary	dept_name	designation	doj	emp_name	project
111	25000	development	associate	2020-06-10 18:30:00.000000+0000	sriram	chat app
121	35000	hr	manager	2020-01-02 18:30:00.000000+0000	shaan	campusx
100	30000	testing	manager	2020-09-10 18:30:00.000000+0000	tanya	canteen app

(3 rows)

```
cqlsh:employee_info> paging off;
```

Disabled Query paging.

```
cqlsh:employee_info> select * from employee_details where emp_id in (111,121,100) order by salary;
```

emp_id	salary	dept_name	designation	doj	emp_name	project
111	25000	development	associate	2020-06-10 18:30:00.000000+0000	sriram	chat app
100	30000	testing	manager	2020-09-10 18:30:00.000000+0000	tanya	canteen app
121	35000	hr	manager	2020-01-02 18:30:00.000000+0000	shaan	campusx

(3 rows)

Date - 26/04/2022

LIBRARY DATABAS E (CASSAND RA)

Question -

Perform the following DB operations using Cassandra.

1. Create a keyspace by name Library
2. Create a column family by name Library-Info with attributes Stud_Id Primary Key, Counter_value of type Counter, Stud_Name, Book-Name, Book-Id, Date_of_issue
3. Insert the values into the table in batch
3. Display the details of the table created and increase the value of the counter
4. Write a query to show that a student with id 112 has taken a book "BDA" 2 times.
5. Export the created column to a csv file
6. Import a given csv dataset from local file system into Cassandra column family

```
cqlsh> create keyspace library_info with replication =  
{'class':'SimpleStrategy','replication_factor':1};
```

```
cqlsh> use library_info;
```

```
cqlsh:library_info> create table library_details(stud_id int,counter_value  
counter,stud_name text,book_name text,date_of_issue timestamp,book_id int,primary  
key(stud_id,stud_name,book_name,date_of_issue,book_id));
```

```
cqlsh:library_info> update library_details set counter_value=counter_value+1 where  
stud_id=111 and stud_name='sam' and book_name='ML' and date_of_issue='2020-11-09' and  
book_id=200;
```

```
cqlsh:library_info> update library_details set counter_value=counter_value+1 where
stud_id=112 and stud_name='shaan' and book_name='BDA' and date_of_issue='2020-01-01'
and book_id=300;
```

```
cqlsh:library_info> update library_details set counter_value=counter_value+1 where
stud_id=113 and stud_name='ayman' and book_name='OOMD' and date_of_issue='2020-06-01'
and book_id=400;
```

```
cqlsh:library_info> select * from library_details;
```

stud_id	stud_name	book_name	date_of_issue	book_id	counter_value
111	sam	ML	2020-11-08 18:30:00.000000+0000	200	1
113	ayman	OOMD	2020-05-31 18:30:00.000000+0000	400	1
112	shaan	BDA	2019-12-31 18:30:00.000000+0000	300	1

(3 rows)

```
cqlsh:library_info> update library_details set counter_value=counter_value+1 where
stud_id=112 and stud_name='shaan' and book_name='BDA' and date_of_issue='2020-01-01'
and book_id=300;
```

```
cqlsh:library_info> select * from library_details where stud_id=112;
```

stud_id	stud_name	book_name	date_of_issue	book_id	counter_value
112	shaan	BDA	2019-12-31 18:30:00.000000+0000	300	2

(1 rows)

```
cqlsh:library_info> copy
```

```
library_details(stud_id,stud_name,book_name,book_id,date_of_issue,counter_value) to  
'E:\sample.csv';
```

Using 3 child processes

Starting copy of library_info.library_details with columns [stud_id, stud_name, book_name,
book_id, date_of_issue, counter_value].

Processed: 3 rows; Rate: 1 rows/s; Avg. rate:

1 rows/s 3 rows exported to 1 files in 3.684 seconds.

```
cqlsh:library_info> truncate library_details;
```

```
cqlsh:library_info> copy
```

```
library_details(stud_id,stud_name,book_name,book_id,date_of_issue,counter_value) from  
'E:\sample.csv';
```

Using 3 child processes

Starting copy of library_info.library_details with columns [stud_id, stud_name, book_name,
book_id, date_of_issue, counter_value].

Processed: 3 rows; Rate: 1 rows/s; Avg. rate: 1

rows/s 3 rows imported from 1 files in 2.602 seconds (0
skipped).

```
cqlsh:library_info> select * from library_details;
```

stud_id	stud_name	book_name	date_of_issue	book_id	counter_value
111	sam	ML	2020-11-08 18:30:00.000000+0000	200	1
113	ayman	OOMD	2020-05-31 18:30:00.000000+0000	400	1
112	shaan	BDA	2019-12-31 18:30:00.000000+0000	300	2

(3 rows)

SCREENSHOTS -

WARNING: pyreadline dependency missing. Install to enable tab completion.

```
cqlsh> create keyspace library_info with replication = {'class':'SimpleStrategy','replication_factor':1};
cqlsh> use library_info;
```

```
ERROR: RequestError from server [code=2100 (invalid query)] message: Unknown definition date_of_issue referenced in library_info
cqlsh:library_info> create table library_details(stud_id int,counter_value counter,stud_name text,book_name text,date_of_issue timestamp,book_id int,primary key(stud_id
,stud_name,book_name,date_of_issue,book_id));
cqlsh:library_info> update library_details set counter_value=counter_value+1 where stud_id=111 and stud_name='sam' and book_name='ML' and date_of_issue='2020-11-09' and
book_id=200;
cqlsh:library_info> update library_details set counter_value=counter_value+1 where stud_id=112 and stud_name='shaan' and book_name='BDA' and date_of_issue='2020-01-01'
and book_id=300;
cqlsh:library_info> update library_details set counter_value=counter_value+1 where stud_id=113 and stud_name='ayman' and book_name='OOMB' and date_of_issue='2020-06-01'
and book_id=400;
cqlsh:library_info> select * from library_details;
```

stud_id	stud_name	book_name	date_of_issue	book_id	counter_value
111	sam	ML	2020-11-08 18:30:00.000000+0000	200	1
113	ayman	OOMB	2020-05-31 18:30:00.000000+0000	400	1
112	shaan	BDA	2019-12-31 18:30:00.000000+0000	300	1

(3 rows)

```
cqlsh:library_info> update library_details set counter_value=counter_value+1 where stud_id=112 and stud_name='shaan' and book_name='BDA' and date_of_issue='2020-01-01'
and book_id=300;
```

```
cqlsh:library_info> select * from library_details where stud_id=112;
```

stud_id	stud_name	book_name	date_of_issue	book_id	counter_value
112	shaan	BDA	2019-12-31 18:30:00.000000+0000	300	2

(1 rows)

```
cqlsh:library_info> copy library_info(stud_id,stud_name,book_name,book_id,date_of_issue,counter_value) to 'E:\sample.csv';
```

```
cqlsh:library_info> copy library_details(stud_id,stud_name,book_name,book_id,date_of_issue,counter_value) to 'E:\sample.csv';
Using 3 child processes
```

```
Starting copy of library_info.library_details with columns [stud_id, stud_name, book_name, book_id, date_of_issue, counter_value].
Processed: 3 rows; Rate:      1 rows/s; Avg. rate:      1 rows/s
3 rows exported to 1 files in 3.684 seconds.
```

```
cqlsh:library_info> truncate library_details;
```

```
cqlsh:library_info> copy library_details(stud_id,stud_name,book_name,book_id,date_of_issue,counter_value) from 'E:\sample.csv';
Using 3 child processes
```

```
Starting copy of library_info.library_details with columns [stud_id, stud_name, book_name, book_id, date_of_issue, counter_value].
Process ImportProcess-6:      1 rows/s; Avg. rate:      1 rows/s
```

```
Processed: 3 rows; Rate:      1 rows/s; Avg. rate:      1 rows/s
3 rows imported from 1 files in 2.602 seconds (0 skipped).
```

```
cqlsh:library_info> select * from library_details;
```

stud_id	stud_name	book_name	date_of_issue	book_id	counter_value
111	sam	ML	2020-11-08 18:30:00.000000+0000	200	1
113	ayman	OOMB	2020-05-31 18:30:00.000000+0000	400	1
112	shaan	BDA	2019-12-31 18:30:00.000000+0000	300	2

(3 rows)

MONGODB SAMPLE

Date - 05/04/2022

Question -

Perform the following DB operations using MongoDB.

1. Create a database “Student” with the following attributes Rollno, Age, ContactNo, Email-Id.
2. Insert appropriate values
3. Write a query to update Email-Id of a student with rollno 10.
4. Replace the student name from “ABC” to “FEM” of rollno 11.
5. Export the created table into local file system
6. Drop the table
7. Import a given csv dataset from the local file system into mongodb collection.

use studentdb

switched to db studentdb

```
db.createCollection("student_details")
```

```
{ "ok" : 1 }
```

```
db.student_details.insert({'name':'abc','rollno':1,'age':19,'contactno':9090909090,'email':'abc@lab.com'})
```

```
WriteResult({ "nInserted" : 1 })
```

```
db.student_details.insert({'name':'mno','rollno':2,'age':20,'contactno':9999900000,'email':'mno@lab.com'})
```

```
WriteResult({ "nInserted" : 1 })
```

```
db.student_details.insert({'name':'xyz','rollno':3,'age':21,'contactno':9999911111,'email':'xyz@lab.com'})
```

```
WriteResult({ "nInserted" : 1 })
```

```
db.student_details.find({})
{ "_id" : ObjectId("60a88f32ffecf7c8abe76775"), "name" : "abc", "rollno" : 1, "age" : 19,
"contactno" : 9090909090, "email" : "abc@lab.com" }
{ "_id" : ObjectId("60a88f7effecf7c8abe76776"), "name" : "mno", "rollno" : 2, "age" : 20,
"contactno" : 9999900000, "email" : "mno@lab.com" }
{ "_id" : ObjectId("60a88f8ffecf7c8abe76777"), "name" : "xyz", "rollno" : 3, "age" :
21, "contactno" : 9999911111, "email" : "xyz@lab.com" }
```

```
db.student_details.update({'rollno':3},{ $set: {'email':'update@lab.com'}})
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
```

```
db.student_details.find({'rollno':3})
{ "_id" : ObjectId("60a88f8ffecf7c8abe76777"), "name" : "xyz", "rollno" : 3, "age" :
21, "contactno" : 9999911111, "email" : "update@lab.com" }
```

```
db.student_details.update({'name':'xyz'}, { $set: {'name':'pqr'}})
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
db.student_details.find({'name':'pqr'})
{ "_id" : ObjectId("60a88f8ffecf7c8abe76777"), "name" : "pqr", "rollno" : 3, "age" :
21, "contactno" : 9999911111, "email" : "update@lab.com" }
```

```
mongoexport --db studentdb --collection student_details --out E:\Desktop\sample.json
2021-05-22T10:43:30.687+0530   connected to: mongodb://localhost/
2021-05-22T10:43:31.026+0530   exported 3 records
```

```
db.getCollection('student_details').drop()
true
```

```
mongoimport --db studentdb --collection student_details --type=json --file=
E:\Desktop\sample.json
2021-05-22T10:46:49.898+0530   connected to: mongodb://localhost/
```


2021-05-22T10:46:50.044+0530 3 document(s) imported successfully. 0 document(s) failed to import.

```
db.student_details.find({})
```

```
{ "_id" : ObjectId("60a88f8ffecf7c8abe76777"), "name" : "pqr", "rollno" : 3, "age" :  
21, "contactno" : 9999911111, "email" : "update@lab.com" }  
{ "_id" : ObjectId("60a88f32ffecf7c8abe76775"), "name" : "abc", "rollno" : 1, "age" : 19,  
"contactno" : 9090909090, "email" : "abc@lab.com" }  
{ "_id" : ObjectId("60a88f7effecf7c8abe76776"), "name" : "mno", "rollno" : 2, "age" : 20,  
"contactno" : 9999900000, "email" : "mno@lab.com" }
```

```
db.student_details.remove({age:{$gt:20}})
```

```
WriteResult({ "nRemoved" : 1 })
```

```
db.student_details.find({})
```

```
{ "_id" : ObjectId("60a88f32ffecf7c8abe76775"), "name" : "abc", "rollno" : 1, "age" : 19,  
"contactno" : 9090909090, "email" : "abc@lab.com" }  
{ "_id" : ObjectId("60a88f7effecf7c8abe76776"), "name" : "mno", "rollno" : 2, "age" : 20,  
"contactno" : 9999900000, "email" : "mno@lab.com" }
```

SCREENSHOTS -

```

> use studentdb
switched to db studentdb
> db.createCollection("student_details")
{ "ok" : 1 }
> db.student_details.insert({'name':'abc','rollno':1,'age':19,'contactno':9090909090,'email':'abc@lab.com'})
WriteResult({ "nInserted" : 1 })
> db.student_details.insert({'name':'mno','rollno':2,'age':20,'contactno':9999900000,'email':'mno@lab.com'})
WriteResult({ "nInserted" : 1 })
> db.student_details.insert({'name':'xyz','rollno':3,'age':21,'contactno':9999911111,'email':'xyz@lab.com'})
WriteResult({ "nInserted" : 1 })
> db.student_details.find({})
{ "_id" : ObjectId("60a88f32ffecf7c8abe76775"), "name" : "abc", "rollno" : 1, "age" : 19, "contactno" : 9090909090, "email" : "abc@lab.com" }
{ "_id" : ObjectId("60a88f7effecf7c8abe76776"), "name" : "mno", "rollno" : 2, "age" : 20, "contactno" : 9999900000, "email" : "mno@lab.com" }
{ "_id" : ObjectId("60a88f8ffecf7c8abe76777"), "name" : "xyz", "rollno" : 3, "age" : 21, "contactno" : 9999911111, "email" : "xyz@lab.com" }
> db.student_details.update({'rollno':3},{ $set:{'email':'update@lab.com'}})
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
> db.student_details.find({'rollno':3})
{ "_id" : ObjectId("60a88f8ffecf7c8abe76777"), "name" : "xyz", "rollno" : 3, "age" : 21, "contactno" : 9999911111, "email" : "update@lab.com" }
> db.student_details.update({'name':'xyz'},{$set:{'name':'pqr'}})
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
> db.student_details.find({'name':'pqr'})
{ "_id" : ObjectId("60a88f8ffecf7c8abe76777"), "name" : "pqr", "rollno" : 3, "age" : 21, "contactno" : 9999911111, "email" : "update@lab.com" }

```

```

C:\Program Files\MongoDB\Server\4.4\bin>mongoexport --db studentdb --collection student_details --out E:\Desktop\sample.json
2021-05-22T10:43:30.687+0530   connected to: mongod://localhost/
2021-05-22T10:43:31.026+0530   exported 3 records

```

```

>
db.getCollection('student_details').drop()
true

```

```

C:\Program Files\MongoDB\Server\4.4\bin>mongoimport --db studentdb --collection student_details --type=json --file= E:\Desktop\sample.json
2021-05-22T10:46:49.898+0530   connected to: mongod://localhost/
2021-05-22T10:46:50.044+0530   3 document(s) imported successfully. 0 document(s) failed to import.

```

```

> db.student_details.find({})
{ "_id" : ObjectId("60a88f32ffecf7c8abe76777"), "name" : "pqr", "rollno" : 3, "age" : 21, "contactno" : 9999911111, "email" : "update@lab.com" }
{ "_id" : ObjectId("60a88f32ffecf7c8abe76775"), "name" : "abc", "rollno" : 1, "age" : 19, "contactno" : 9090909090, "email" : "abc@lab.com" }
{ "_id" : ObjectId("60a88f7effecf7c8abe76776"), "name" : "mno", "rollno" : 2, "age" : 20, "contactno" : 9999900000, "email" : "mno@lab.com" }
> db.student_details.remove({'age':{'$gt':20}})
WriteResult({ "nRemoved" : 1 })
> db.student_details.find({})
{ "_id" : ObjectId("60a88f32ffecf7c8abe76775"), "name" : "abc", "rollno" : 1, "age" : 19, "contactno" : 9090909090, "email" : "abc@lab.com" }
{ "_id" : ObjectId("60a88f7effecf7c8abe76776"), "name" : "mno", "rollno" : 2, "age" : 20, "contactno" : 9999900000, "email" : "mno@lab.com" }

```

SCREENSHOT OF HADOOP INSTALLATION

Date - 12/04/2022

```
C:\Users\Admin>hadoop version
Hadoop 3.1.0
Source code repository https://github.com/apache/hadoop -r 16b70619a24cdcf5d3b0fcf4b58ca77238ccbe6d
Compiled by centos on 2018-03-30T00:00Z
Compiled with protoc 2.5.0
From source with checksum 14182d20c972b3e2105580a1ad6990
This command was run using /C:/hadoop_new/share/hadoop/common/hadoop-common-3.1.0.jar

C:\Users\Admin>cd c:\hadoop_new\sbin

c:\hadoop_new\sbin>start-all.cmd
This script is Deprecated. Instead use start-dfs.cmd and start-yarn.cmd
starting yarn daemons
```

HADOOP SAMPLE

Date - 19/04/2022

Execution of HDFS Commands for interaction with Hadoop Environment. (Minimum 10 commands to be executed)

```
c:\hadoop_new\sbin>hdfs dfs -mkdir /temp
c:\hadoop_new\sbin>hdfs dfs -copyFromLocal E:\Desktop\sample.txt \temp
c:\hadoop_new\sbin>hdfs dfs -ls \temp
```

Found 1 items

```
-rw-r--r-- 1 Admin supergroup      11 2021-06-11 21:12 /temp/sample.txt
```

```
c:\hadoop_new\sbin>hdfs dfs -cat \temp\sample.txt
```

hello world

```
c:\hadoop_new\sbin>hdfs dfs -get \temp\sample.txt E:\Desktop\temp
```

```
c:\hadoop_new\sbin>hdfs dfs -put E:\Desktop\temp \temp
```

```
c:\hadoop_new\sbin>hdfs dfs -ls \temp
```

Found 2 items

```
-rw-r--r-- 1 Admin supergroup      11 2021-06-11 21:12 /temp/sample.txt
drwxr-xr-x - Admin supergroup      0 2021-06-11 21:15 /temp/temp
```

```
c:\hadoop_new\sbin>hdfs dfs -mv \lab1 \temp
```

```
c:\hadoop_new\sbin>hdfs dfs -ls \temp
```

Found 3 items

```
drwxr-xr-x - Admin supergroup      0 2021-04-19 15:07 /temp/lab1
```

```
-rw-r--r-- 1 Admin supergroup      11 2021-06-11 21:12 /temp/sample.txt
drwxr-xr-x - Admin supergroup      0 2021-06-11 21:15 /temp/temp
```

```
c:\hadoop_new\sbin>hdfs dfs -rm /temp/sample.txt
```

```
Deleted /temp/sample.txt
```

```
c:\hadoop_new\sbin>hdfs dfs -ls /temp
```

```
Found 2 items
```

```
drwxr-xr-x - Admin supergroup      0 2021-04-19 15:07 /temp/lab1
drwxr-xr-x - Admin supergroup      0 2021-06-11 21:15 /temp/temp
c:\hadoop_new\sbin>hdfs dfs -copyFromLocal E:\Desktop\sample.txt /temp
c:\hadoop_new\sbin>hdfs dfs -ls /temp
```

```
Found 3 items
```

```
drwxr-xr-x - Admin supergroup      0 2021-04-19 15:07 /temp/lab1
-rw-r--r-- 1 Admin supergroup      11 2021-06-11 21:17 /temp/sample.txt
drwxr-xr-x - Admin supergroup      0 2021-06-11 21:15 /temp/temp
```

```
c:\hadoop_new\sbin>hdfs dfs -copyToLocal /temp/sample.txt
```

```
E:\Desktop\sample.txt SCREENSHOTS -
```

```

: ',hadoop_nen',sbin>hdfs dfs -m<dir Temp
: ',hadoop_nen',sbin>hdfs dfs -:op; FromLo:al E:',Des<:op',sample.:x: ',:emp
: ',hadoop_nen',sbin>hdfs dfs -ls
,:emp Found 1 i:ems
-nn-n--n-- 1 Admin supengnoup 11 2021-0o-11 21:12 ,':emp/sample.:x:
: ',hadoop_nen',sbin>hdfs dfs -ca:
',:emp\sample.:x: hello world
: ',hadoop_nen',sbin>hdfs dfs -ge: ',:emp\sample.:x: E:',Des<:op',:emp
: ',hadoop_nen',sbin>hdfs dfs -pu: E:',Des<:op',:emp ',:emp
: ',hadoop_nen',sbin>hdfs dfs -ls ,:emp 11 2021-0o-11 21:12 ,':emp/sample.:x:
Found 2 items 0 2021-0o-11 21:15 ,':emp;:emp
-nn-n--n-- 1 Admin supengnoup
drnxr-xr-x - Admin supergroup

```

```

: ',hadoop_nen',sbin>hdfs dfs -mv ',labl ',:emp
: ',hadoop_nen',sbin>hdfs dfs -ls
',:emp Found 3 i:ems
druxr-xr-x - Admin supergroup 0 2021-04-19 15:07 ;:emp/labl
-ru-r--r-- 1 Admin supergroup 11 2021-0o-11 21:12
druxr-xr-x - Admin supergroup ;:emp/sample.:x:
0 2021-0o-11 21:15 ;:emp;:emp

```

```

: ',hadoop_nen',sbin>hdfs dfs -rm
;:emp/sample.:x: Deleted
;:emp/sample.:x:

```

```

: ',hadoop_nen',sbin>hdfs dfs 0 2021-04-19 15:07 /-emp/labl
-ls ',:emp Found 2 i:ems 0 2021-0o-11 21:15 ;:emp;:emp
dnuxn-xn-x Admin
supengnoup dnuxn-xn-x -
Admin supengnoup

```

```

: ',hadoop_neu',sbin>hdfs dfs -copyFnomLo:al E: ,Des<-op ,sample.-x- ,:emp

```

```

: ',hadoop_nen',sbin>hdfs dfs -ls 0 2021-04-19 15:07 ;:emp/labl
',:emp Found 3 i:ems 11 2021-0o-11 21:17 ;:emp/sample.:x:
druxr-xr-x - Admin supergroup 0 2021-0o-11 21:15 ;:emp;:emp
-ru-r--r-- 1 Admin supergroup
druxr-xr-x - Admin supergroup

```

```

: ',hadoop_nen',sbin>hdfs dfs -:opy oLo:al ',:emp',sample.:x:
E:',Des<:op',sample.:x:

```

MAPREDUCE TEMPERATURE

Date - 10/05/2022

For the given file, Create a Map Reduce program to

a) Find the average temperature for each year from the NCDC data set.

```
// AverageDriver.java
package temperature;

import org.apache.hadoop.io.*;
import org.apache.hadoop.fs.*;
import org.apache.hadoop.mapreduce.*;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class AverageDriver
{
    public static void main (String[] args) throws Exception
    {
        if (args.length != 2)
        {
            System.err.println("Please Enter the input and output
            parameters"); System.exit(-1);
        }
        Job job = new Job();
        job.setJarByClass(AverageDriver.class);
        job.setJobName("Max temperature");
        FileInputFormat.addInputPath(job,new Path(args[0]));
        FileOutputFormat.setOutputPath(job,new Path (args[1]));

        job.setMapperClass(AverageMapper.class);
```

```

        job.setReducerClass(AverageReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        System.exit(job.waitForCompletion(true)?0:1);
    }
}

```

//AverageMapper.java

```
package temperature;
```

```
import org.apache.hadoop.io.*;
```

```
import org.apache.hadoop.mapreduce.*;
```

```
import java.io.IOException;
```

```
public class AverageMapper extends Mapper <LongWritable, Text, Text, IntWritable>
```

```
{
```

```
    public static final int MISSING = 9999;
```

```
    public void map(LongWritable key, Text value, Context context) throws
```

```
        IOException, InterruptedException
```

```
{
```

```
        String line = value.toString();
```

```
        String year = line.substring(15,19);
```

```
        int temperature;
```

```
        if (line.charAt(87)=='+')

```

```
            temperature = Integer.parseInt(line.substring(88, 92));
```

```
        else

```

```
            temperature = Integer.parseInt(line.substring(87, 92));
```

```
        String quality = line.substring(92, 93);
```

```
        if(temperature != MISSING && quality.matches("[01459]"))

```

```
            context.write(new Text(year),new
```

```
                IntWritable(temperature));

```



```

    }
}

//AverageReducer.java
package temperature;

import
org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.*;
import java.io.IOException;

public class AverageReducer extends Reducer <Text, IntWritable,Text, IntWritable>
{
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
throws IOException,InterruptedException
    {
        int max_temp = 0;
        int count = 0;
        for (IntWritable value : values)
        {
            max_temp += value.get();
            count+=1;
        }
        context.write(key, new IntWritable(max_temp/count));
    }
}

```

SCREENSHOT -

```
c:\hadoop_new\sbin>hdfs dfs -cat /tempAverageOutput/part-r-00000
1901      46
1949      94
1950       3
```

b) Find the mean max temperature for every month.

```
//TempDriver.java
```

```
package temperatureMax;
```

```
import org.apache.hadoop.io.*;
```

```
import org.apache.hadoop.fs.*;
```

```
import org.apache.hadoop.mapreduce.*;
```

```
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
```

```
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

```
public class TempDriver
```

```
{
```

```
    public static void main (String[] args) throws Exception
```

```
    {
```

```
        if (args.length != 2)
```

```
        {
```

```
            System.err.println("Please Enter the input and output  
parameters"); System.exit(-1);
```

```
        }
```

```
        Job job = new Job();
```

```
        job.setJarByClass(TempDriver.class);
```

```
        job.setJobName("Max temperature");
```

```
        FileInputFormat.addInputPath(job,new Path(args[0]));
```

```
        FileOutputFormat.setOutputPath(job,new Path (args[1]));
```

```
        job.setMapperClass(TempMapper.class);
```

```

        job.setReducerClass(TempReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        System.exit(job.waitForCompletion(true)?0:1);
    }
}

//TempMapper.java
package temperatureMax;

import org.apache.hadoop.io.*;
import org.apache.hadoop.mapreduce.*;
import java.io.IOException;

public class TempMapper extends Mapper <LongWritable, Text, Text, IntWritable>
{
    public static final int MISSING = 9999;

    public void map(LongWritable key, Text value, Context context) throws
    IOException, InterruptedException
    {
        String line = value.toString();
        String month = line.substring(19,21);
        int temperature;
        if (line.charAt(87)=='+')
            temperature = Integer.parseInt(line.substring(88, 92));
        else
            temperature = Integer.parseInt(line.substring(87, 92));
        String quality = line.substring(92, 93);
        if(temperature != MISSING && quality.matches("[01459]"))
            context.write(new Text(month),new
            IntWritable(temperature));
    }
}

```

```

    }
}

//TempReducer.java
package temperatureMax;

import org.apache.hadoop.io.*;
import org.apache.hadoop.mapreduce.*;
import java.io.IOException;

public class TempMapper extends Mapper <LongWritable, Text, Text, IntWritable>
{
    public static final int MISSING = 9999;

    public void map(LongWritable key, Text value, Context context) throws
    IOException, InterruptedException
    {
        String line = value.toString();
        String month = line.substring(19,21);
        int temperature;
        if (line.charAt(87)=='+')
            temperature = Integer.parseInt(line.substring(88, 92));
        else
            temperature = Integer.parseInt(line.substring(87, 92));
        String quality = line.substring(92, 93);
        if(temperature != MISSING && quality.matches("[01459]"))
            context.write(new Text(month),new
            IntWritable(temperature));
    }
}

```

SCREENSHOT -

317

```
c:\hadoop_new\sbin>hdfs dfs -cat /tempMaxOutput/part-r-00000
```

02

03

111

04

124

05

125

06

156

10

89

11

117

MAPREDUCE TOPN

Date - 03/05/2022

For a given Text file, create a Map Reduce program to sort the content in an alphabetic order listing only top 'n' maximum occurrence of words.

```
// TopN.java
package sortWords;

import
org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import
org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.util.GenericOptionsParser;
import utils.MiscUtils;

import java.io.IOException;
import java.util.*;

public class TopN {

    public static void main(String[] args) throws Exception
    { Configuration conf = new Configuration();
      String[] otherArgs = new GenericOptionsParser(conf, args).getRemainingArgs();
      if (otherArgs.length != 2) {
```

```
System.err.println("Usage: TopN <in> <out>");
```

```

        System.exit(2);
    }

    Job job = Job.getInstance(conf);
    job.setJobName("Top N");
    job.setJarByClass(TopN.class);
    job.setMapperClass(TopNMapper.class);
    //job.setCombinerClass(TopNReducer.class);
    job.setReducerClass(TopNReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(otherArgs[0]));
    FileOutputFormat.setOutputPath(job, new
    Path(otherArgs[1])); System.exit(job.waitForCompletion(true)
    ? 0 : 1);
}

/**
 * The mapper reads one line at the time, splits it into an array of single words and emits every
 * word to the reducers with the value of 1.
 */

public static class TopNMapper extends Mapper<Object, Text, Text, IntWritable> {

    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();
    private String tokens = "[_!$%&^`=\\[\\]\\\\*\\/\\\\\\\\,;\\.\\-:()?!\\\"'"]";

    @Override
    public void map(Object key, Text value, Context context) throws
    IOException, InterruptedException {
        String cleanLine = value.toString().toLowerCase().replaceAll(tokens, " ");
        StringTokenizer itr = new StringTokenizer(cleanLine);
        while (itr.hasMoreTokens()) {

```



```

        word.set(itr.next().trim());
        context.write(word, one);
    }
}
}

```

```
/**
```

```

 * The reducer retrieves every word and puts it into a Map: if the word already exists in the
 * map, increments its value, otherwise sets it to 1.
 */

```

```
*/
```

```
public static class TopNReducer extends Reducer<Text, IntWritable, Text, IntWritable>
```

```
{ private Map<Text, IntWritable> countMap = new HashMap<>();
```

```
@Override
```

```
public void reduce(Text key, Iterable<IntWritable> values, Context context)
```

```
throws IOException, InterruptedException {
```

```
    // computes the number of occurrences of a single word
```

```
    int sum = 0;
```

```
    for (IntWritable val : values)
```

```
    { sum += val.get();
```

```
    }
```

```
    // puts the number of occurrences of this word into the map.
```

```
    // We need to create another Text object because the Text instance
```

```
    // we receive is the same for all the words
```

```
    countMap.put(new Text(key), new
```

```
    IntWritable(sum));
```

```
}
```

```
@Override
```

```

protected void cleanup(Context context) throws IOException, InterruptedException {

    Map<Text, IntWritable> sortedMap = MiscUtils.sortByValues(countMap);

    int counter = 0;
    for (Text key : sortedMap.keySet()) {
        if (counter++ == 3) {
            break;
        }
        context.write(key, sortedMap.get(key));
    }
}

/**
 * The combiner retrieves every word and puts it into a Map: if the word already exists in the
 * map, increments its value, otherwise sets it to 1.
 */
public static class TopNCombiner extends Reducer<Text, IntWritable, Text, IntWritable> {

    @Override
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
throws IOException, InterruptedException {

        // computes the number of occurrences of a single word
        int sum = 0;
        for (IntWritable val : values)
            { sum += val.get();
        }
        context.write(key, new IntWritable(sum));
    }
}

```

```

    }
}

// MiscUtils.java
package utils;

import java.util.*;

public class MiscUtils

{

    /**
     * sorts the map by values. Taken from:
     * http://javarevisited.blogspot.it/2012/12/how-to-sort-hashmap-java-by-key-and-value.html
     */
    public static <K extends Comparable, V extends Comparable> Map<K,
V> sortByValues(Map<K, V> map) {
        List<Map.Entry<K, V>> entries = new LinkedList<Map.Entry<K,

V>>(map.entrySet()); Collections.sort(entries, new Comparator<Map.Entry<K, V>>() {

        @Override
        public int compare(Map.Entry<K, V> o1, Map.Entry<K, V> o2)
        { return o2.getValue().compareTo(o1.getValue());
        }
    });

    //LinkedHashMap will keep the keys in the order they are inserted
    //which is currently sorted on natural ordering
    Map<K, V> sortedMap = new LinkedHashMap<K,

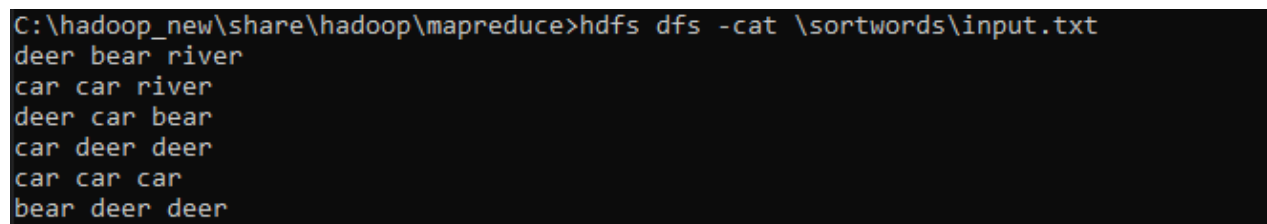
V>(); for (Map.Entry<K, V> entry : entries) {

```

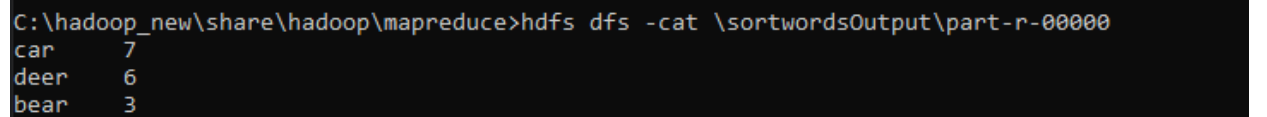
```
        sortedMap.put(entry.getKey(), entry.getValue());
    }

    return sortedMap;
}
```

SCREENSHOTS -



```
C:\hadoop_new\share\hadoop\mapreduce>hdfs dfs -cat \sortwords\input.txt
deer bear river
car car river
deer car bear
car deer deer
car car car
bear deer deer
```



```
C:\hadoop_new\share\hadoop\mapreduce>hdfs dfs -cat \sortwordsOutput\part-r-00000
car      7
deer     6
bear     3
```

MAPREDUCE JOIN

Date - 31/06/2022

Create a Hadoop Map Reduce program to combine information from the users file along with Information from the posts file by using the concept of join and display user_id, Reputation and Score.

```
// JoinDriver.java
```

```
import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;
import org.apache.hadoop.mapred.libMultipleInputs;
import org.apache.hadoop.util.*;
```

```
public class JoinDriver extends Configured implements Tool {
```

```
    public static class KeyPartitioner implements Partitioner<TextPair, Text> {
        @Override
        public void configure(JobConf job) {}
```

```
numPartitions;
```

```
        @Override
        public int
            getPartition(TextPair key,
                Text value, int
                numPartitions)
            { return
                (key.getFirst().
                hashCode() &
                Integer.MAX_
                VALUE) %
```

```
    }  
}
```

```
@Override  
public int run(String[] args) throws Exception {
```

```

        if (args.length != 3) {
            System.out.println("Usage: <Department Emp Strength input>
<Department Name input> <output>");
            return -1;
        }

```

```

input");

```

```

Job
Conf
conf
=
new
Job
Conf
f(
getC
onf()
,
getC
lass()
);
conf.se
tJobN
ame("
Join
'De
partme
nt
Emp
p
Str
eng

```

P
a
t
h
A
I
n
p
u
t
P
a
t
h
=
n
e
w
P
a
t
h
(
a
r

th
inp
ut'
wit
h
'De
par
tme
nt
Na
me

g
s
[
0
]
)
;
P
a
t
h
B
I
n
p
u
t
P
a
t
h
=
n
e
w
P
a
t
h
(
a
r
g
s


```
[
1
]
)
;
P
a
t
h
o
u
t
p
u
t
P
a
t
h
=
n
e
w
P
a
t
h
(
a
r
g
s
[
2
]
)
;
```

```
MultipleInputs.addInputPath(conf,
Posts.class); User.class);
```

```
AInputPath,
```

```
TextInputFormat.class,
```

```
MultipleInputs.addInputPath(conf,
```

```
BInputPath,
```

```
TextInputFormat.class,
```

```
FileOutputFormat.setOutputPath(conf, outputPath);
```

```
conf.setPartitionerClass(KeyPartitioner.class);
```

```
conf.setOutputValueGroupingComparator(TextPair.FirstComparator.class);
```

```
conf.setMapOutputKeyClass(TextPair.class);  
conf.setReducerClass(JoinReducer.class);  
conf.setOutputKeyClass(Text.class);
```

```

        JobClient.runJob(conf);

        return 0;
    }

    public static void main(String[] args) throws Exception {

        int exitCode = ToolRunner.run(new JoinDriver(), args);
        System.exit(exitCode);
    }
}

// JoinReducer.java
import java.io.IOException;
import java.util.Iterator;

import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;

public class JoinReducer extends MapReduceBase implements Reducer<TextPair, Text,
Text, Text> {

    @Override
    public void reduce (TextPair key, Iterator<Text> values, OutputCollector<Text,
Text> output, Reporter reporter)
        throws IOException
    {

        Text nodeId = new
        Text(values.next()); while
        (values.hasNext()) {
            Text node = values.next();

```

```

        Text outValue = new Text(nodeId.toString() + "\t\t" +
        node.toString()); output.collect(key.getFirst(), outValue);
    }
}

```

// User.java

```

import java.io.IOException;
import java.util.Iterator;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.FSDataInputStream;
import org.apache.hadoop.fs.FSDataOutputStream;
import org.apache.hadoop.fs.FileSystem;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;

import org.apache.hadoop.io.IntWritable;

public class User extends MapReduceBase implements Mapper<LongWritable, Text,
TextPair, Text> {

    @Override
    public void map(LongWritable key, Text value, OutputCollector<TextPair, Text>
output, Reporter reporter)
        throws IOException
    {

        String valueString = value.toString();
        String[] SingleNodeData = valueString.split("\t");
    }
}

```

```
        output.collect(new TextPair(SingleNodeData[0], "1"),
new Text(SingleNodeData[1]));
    }
}
```

//Posts.java

```
import java.io.IOException;
```

```
import org.apache.hadoop.io.*;
```

```
import org.apache.hadoop.mapred.*;
```

```
public class Posts extends MapReduceBase implements Mapper<LongWritable, Text,
TextPair, Text> {
```

```
    @Override
```

```
    public void map(LongWritable key, Text value, OutputCollector<TextPair, Text>
output, Reporter reporter)
```

```
        throws IOException
```

```
    {
```

```
        String valueString = value.toString();
```

```
        String[] SingleNodeData = valueString.split("\t");
```

```
        output.collect(new TextPair(SingleNodeData[3], "0"),
```

```
        new
```

```
Text(SingleNodeData[9]));
```

```
    }
```

```
}
```

// TextPair.java

```
import java.io.*;
```

```
import org.apache.hadoop.io.*;
```

```
public class TextPair implements WritableComparable<TextPair> {
```

```
    private Text first;
```

```
    private Text second;
```

```
    public TextPair() {
```

```
        set(new Text(), new Text());
```

```
    }
```

```
    public TextPair(String first, String second)
```

```
    { set(new Text(first), new Text(second));
```

```
    }
```

```
    public TextPair(Text first, Text second)
```

```
    { set(first, second);
```

```
    }
```

```
    public void set(Text first, Text second)
```

```
    { this.first = first;
```

```
      this.second = second;
```

```
    }
```

```
    public Text getFirst() {
```

```
        return first;
```

```
    }
```

```
    public Text getSecond()
```

```
    { return second;
```

```
    }
```

```
    @Override
```

```
public void write(DataOutput out) throws IOException {  
    first.write(out);  
    second.write(out);  
}
```

@Override

```
public void readFields(DataInput in) throws IOException {  
    first.readFields(in);  
    second.readFields(in);  
}
```

@Override

```
public int hashCode() {  
    return first.hashCode() * 163 + second.hashCode();  
}
```

@Override

```
public boolean equals(Object o) {  
    if (o instanceof TextPair) {  
        TextPair tp = (TextPair) o;  
        return first.equals(tp.first) && second.equals(tp.second);  
    }  
    return false;  
}
```

@Override

```
public String toString() {  
    return first + "\t" +  
        second;  
}
```

@Override

```

public int compareTo(TextPair tp) {
    int cmp = first.compareTo(tp.first);
    if (cmp != 0) {
        return cmp;
    }
    return second.compareTo(tp.second);
}
// ^^ TextPair

// vv TextPairComparator
public static class Comparator extends WritableComparator {
    private static final Text.Comparator TEXT_COMPARATOR = new Text.Comparator();

    public Comparator() {

        super(TextPair.class);
    }

    @Override
    public int compare(byte[] b1, int s1, int l1,
        byte[] b2, int s2, int l2) {

        try {
            int firstL1 = WritableUtils.decodeVIntSize(b1[s1]) + readVInt(b1,
                s1); int firstL2 = WritableUtils.decodeVIntSize(b2[s2]) +
                readVInt(b2, s2);
            int cmp = TEXT_COMPARATOR.compare(b1, s1, firstL1, b2, s2, firstL2);
            if (cmp != 0) {
                return cmp;
            }
            return TEXT_COMPARATOR.compare(b1, s1 + firstL1, l1 -
                firstL1, b2, s2 + firstL2, l2 - firstL2);
        }
    }
}

```



```

    } catch (IOException e) {
        throw new IllegalArgumentException(e);
    }
}

```

```

static {
    WritableComparator.define(TextPair.class, new Comparator());
}

public static class FirstComparator extends WritableComparator {
    private static final Text.Comparator TEXT_COMPARATOR = new
    Text.Comparator(); public FirstComparator() {
        super(TextPair.class);
    }
}

```

@Override

```

public int compare(byte[] b1, int s1, int l1,
    byte[] b2, int s2, int l2) {

```

```

    try {
        int firstL1 = WritableUtils.decodeVIntSize(b1[s1]) + readVInt(b1, s1);
        int firstL2 = WritableUtils.decodeVIntSize(b2[s2]) + readVInt(b2, s2);
        return TEXT_COMPARATOR.compare(b1, s1, firstL1, b2, s2,
            firstL2);
    } catch (IOException e) {
        throw new IllegalArgumentException(e);
    }
}

```

@Override

```

public int compare(WritableComparable a, WritableComparable b)
{
    if (a instanceof TextPair && b instanceof TextPair) {
        return ((TextPair) a).first.compareTo(((TextPair) b).first);
    }
    return super.compare(a, b);
}
}
}
}

```

SCREENSHOTS -

```

c:\hadoop_new\share\hadoop\mapreduce>hdfs dfs -cat \posts\sampleposts.tsv
"2312" "Feedback on Audio Quality" "cs101 production audio" "100005361" "<p>We are looking for feedback on the audio in our videos. Tell us what you thi
nk and try to be as <em>specific</em> as possible.</p>" "question"
"\N" "\N" "2012-02-23 00:28:02.321344+00" "2" "" "\N" "201398145" "2014-01-14 17:18:35.613939+00" "2960" "\N" "\N" "524" "f"
"2014856" "" "cs101" "100022094" "<p>I also would like to know the answer to this question. An 'open exam' sounds great, but on the other hand it
also seems pretty easy to cheat now: solutions have been posted and anybody only interested in a certificate wouldn't have much of a problem getting the highest distin
ction. So where is the catch??</p>" "answer" "2014706" "2014706" "2012-07-01 10:32:36.302782+00" "0" "" "\N"
"100022094" "2012-07-01 10:32:36.302782+00" "2020501" "\N" "\N" "0" "f"
"2004004" "" "cs101" "100018705" "<p>But then why even the new variable q? Why not just modify the variable p?</p>" "comment" "2003997"
"2003993" "2012-05-03 21:07:52.028935+00" "2" ""
"\N" "100018705" "2012-05-03 21:07:52.028935+00" "2005150" "\N" "\N" "0" "f"

c:\hadoop_new\share\hadoop\mapreduce>hdfs dfs -cat \users\sampleusers.tsv
"100006402" "18" "0" "0" "0"
"100022094" "6354" "4" "12" "50"
"100018705" "76" "0" "3" "4"
"100005361" "36134" "73" "220" "333"

```

```

c:\hadoop_new\share\hadoop\mapreduce>hdfs dfs -cat \joinOutput\part-00000
"100005361" "2" "36134"
"100018705" "2" "76"
"100022094" "0" "6354"

```

SCALA INSTALLATION SCREENSHOT

```
sam@ubuntu:~$ start-master.sh
starting org.apache.spark.deploy.master.Master, logging to /opt/spark/logs/spark-sam-org.apache.spark.deploy.mas
ter.Master-1-ubuntu.out
sam@ubuntu:~$ start-slave.sh spark://ubuntu:7077
This script is deprecated, use start-worker.sh
starting org.apache.spark.deploy.worker.Worker, logging to /opt/spark/logs/spark-sam-org.apache.spark.deploy.wor
ker.Worker-1-ubuntu.out
```

```
sam@ubuntu:~$ spark-shell
21/06/13 07:19:08 WARN Utils: Your hostname, ubuntu resolves to a loopback address: 127.0.1.1; using 192.168.18.
128 instead (on interface ens33)
21/06/13 07:19:08 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.spark.unsafe.Platform (file:/opt/spark/jars/spark-unsafe_2.12-3
.1.1.jar) to constructor java.nio.DirectByteBuffer(long,int)
WARNING: Please consider reporting this to the maintainers of org.apache.spark.unsafe.Platform
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
21/06/13 07:19:10 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin
-java classes where applicable
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Spark context Web UI available at http://192.168.18.128:4040
Spark context available as 'sc' (master = local[*], app id = local-1623593969342).
Spark session available as 'spark'.
Welcome to

  ____      _
 / ___|  _ \| | | |
 \___ \ |_) | |_| |
  ___) |  __/| | | |
 |____|_| \_||_|_|_|

 version 3.1.1

Using Scala version 2.12.10 (OpenJDK 64-Bit Server VM, Java 11.0.11)
Type in expressions to have them evaluated.
Type :help for more information.

scala>
```

SCALA WORDCOUNT

Date - 07/06/2022

```
// scala shell
```

```
scala> val textfile = sc.textFile("/home/sam/Desktop/abc.txt")
```

```
textfile: org.apache.spark.rdd.RDD[String] = /home/sam/Desktop/abc.txt MapPartitionsRDD[1]
```

```
at textFile at <console>:24
```

```
scala> val counts = textfile.flatMap(line => line.split(" ")).map(word =>
```

```
(word,1)).reduceByKey(+)
```

```
counts: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[4] at reduceByKey at
```

```
<console>:25
```

```
scala> import scala.collection.immutable.ListMap
```

```
import scala.collection.immutable.ListMap
```

```
scala> val sorted = ListMap(counts.collect.sortWith(_. _2>.2):*)
```

```
scala> println(sorted)
```

```
ListMap(car -> 7, deer -> 5, bear -> 3, river -> 3, -> 1)
```

```
scala> for((k,v)<-sorted)
```

```
| {
```

```
| if(v>4)
```

```
| {
```

```
| println(k+"-"+v)
```

```
| }
```

```
| }
```

```
car-7
```

```
deer-5
```

```

val textfile = sc.textFile("phone/sam/Desktop/abc.txt")
// /home/sam/Desktop/abc.txt MapPartitionsRDD[8] at textFile
// at <console>:25

let:25

val counts = textfile.flatMap(line => line.split(" ").map(word => (word,1)).reduceByKey(_+_))
// ShuffledRDD[11] at reduceByKey at <console>:26

import
scala.collection.immutable.ListMap import
scala.collection.immutable.ListMap

val sorted = ListMap(counts.collect.sortWith(_._2 > _._2):_*)
// ListMap(hello -> 3, apple -> 2, unicorn -> 1,
// world -> 1)

println(sorted)
// ListMap(hello -> 3, apple -> 2, unicorn -> 1, world -> 1)

```

```

for ((k, v) <- sorted)
  if (v > 2)

```

```

    println(k + " - " + v)

```

/wc.txt")

```

val counts = textfile.flatMap(line => line.split(" ").map(word => (word, 1)).reduceByKey(_+_))

```

```

// ListMap(hello -> 3, apple -> 2, unicorn -> 1, world -> 1)

```

```

val sorted = ListMap(counts.collect.sortWith(_._2 > _._2):_*) // sort in descending order based on values

```

```

println(sorted)

```

```

for ((k, v) <- sorted)

```

```

{

```

```

  if (v > 4)

```

```

  {

```

```

    print(k + ",")

```

```

    print(v)

```

```

    println()

```

```

  }

```

```

}

```